

# Homework 3: Augmented Reality with Planar Homographies

For each question please refer to the handout for more details.

Programming questions begin at Q2. Remember to run all cells and save the notebook to your local machine as a pdf for gradescope submission.

## Collaborators

List your collaborators for all questions here:

### Q1 Preliminaries

#### Q1.1 The Direct Linear Transform

##### Q1.1.1 (3 points)

How many degrees of freedom does  $\mathbf{h}$  have?

$\mathbf{h}$  has 8 degrees of freedom. Although  $\mathbf{h}$  is a 3x3 matrix,  $\mathbf{h}$  is defined up to a scale factor. This makes it possible for the 9th value to have a value of 1, giving exactly 8 degrees of freedom.

##### Q1.1.2 (2 points)

How many point pairs are required to solve  $\mathbf{h}$ ?

To solve  $\mathbf{h}$ , we would require 4 point pairs. Each point pair would give us 2 equations. With 4 points, we would have 8 equations for our 8 unknowns in  $\mathbf{A}$ , making it sufficient to solve  $\mathbf{h}$ .

##### Q1.1.3 (5 points)

Derive  $\mathbf{A}_i$

##### Q1.1.4 (5 points)

What will be the trivial solution for  $\mathbf{h}$ ? Is the matrix  $\mathbf{A}$  full rank? Why/Why not? What impact will it have on the singular values (i.e. eigenvalues of  $\mathbf{A}^T \mathbf{A}$ )?

The trivial solution for  $\mathbf{A}\mathbf{h}=0$  would be a matrix  $\mathbf{h}$  of all zeros. This would essentially represent no transformation. It isn't full rank because it represents the correspondence between points in two different planes, leading to linearly dependent rows. Some singular values will be zero, indicating solutions with minimal effect.

#### Q1.2 Homography Theory Questions

##### Q1.2.1 (5 points)

Prove that there exists a homography  $\mathbf{H}$  that satisfies  $\mathbf{x}_1 = \mathbf{H}\mathbf{x}_2$ , given two cameras separated by a pure rotation.

YOUR ANSWER HERE...

##### Q1.2.2 (5 points):

Show that  $\mathbf{H}^2$  is the homography corresponding to a rotation of  $2\theta$ .

This is equivalent to  $\mathbf{H}\mathbf{H}$

$$\mathbf{H}^2 = (\mathbf{K} \mathbf{R}(\theta) \mathbf{K}^{-1})^2$$

$$= \mathbf{K} \mathbf{R}(\theta) \mathbf{K}^{-1} \mathbf{K} \mathbf{R}(\theta) \mathbf{K}^{-1}$$

$$= \mathbf{K} \mathbf{R}(\theta) \mathbf{R}(\theta) \mathbf{K}^{-1}$$

$$= \mathbf{K} \mathbf{R}(2\theta) \mathbf{K}^{-1}$$

This is because upon multiplication, the terms correspond to  $\sin(2\theta)$  and  $\cos(2\theta)$

$$[[\cos^2(\theta) - \sin^2(\theta), -2\sin(\theta)\cos(\theta)], [2\sin(\theta)\cos(\theta), -\sin^2(\theta)+\cos^2(\theta)]]$$

which is equivalent to:

$$[[\cos(2\theta), \sin(2\theta)], [\sin(2\theta), \cos(2\theta)]]$$

$$= \mathbf{R}(2\theta)$$

#### Initialization

Run the following code to import the modules you'll need.

```

1 import os
2 import numpy as np
3 import cv2
4 import skimage.color
5 import pickle
6 from matplotlib import pyplot as plt
7 import scipy
8 from skimage.util import montage
9 import time
10 from skimage.io import imread
11
12 PATCHWIDTH = 9
13
14 def read_pickle(path):
15     with open(path, "rb") as f:
16         return pickle.load(f)
17
18 def write_pickle(path, data):
19     with open(path, "wb") as f:
20         pickle.dump(data, f)
21
22 def briefMatch(desc1, desc2, ratio):
23
24     matches = skimage.feature.match_descriptors(desc1, desc2,
25                                                'hamming',
26                                                cross_check=True,
27                                                max_ratio=ratio)
28     return matches
29
30 def plotMatches(img1, img2, matches, locs1, locs2):
31
32     fig, ax = plt.subplots(nrows=1, ncols=1)
33     """
34     print("cvtColor")
35     print(img1.shape)
36     print(img2.shape)
37
38     if len(img1.shape) > 2:
39         I1_gray = skimage.color.rgb2gray(img1)
40     else:
41         I1_gray = img1
42
43     if len(img2.shape) > 2:
44         I2_gray = skimage.color.rgb2gray(img2)
45     else:
46         I2_gray = img2
47     """
48
49     img1 = cv2.cvtColor(img1, cv2.COLOR_BGR2GRAY)
50     img2 = cv2.cvtColor(img2, cv2.COLOR_BGR2GRAY)
51     plt.axis('off')
52     skimage.feature.plot_matches(ax, img1, img2, locs1, locs2,
53                                matches, matches_color='r', only_matches=True)
54     plt.show()
55     return
56
57 def makeTestPattern(patchWidth, nbits):
58
59     np.random.seed(0)
60     compareX = patchWidth*patchWidth * np.random.random((nbits,1))
61     compareX = np.floor(compareX).astype(int)
62     np.random.seed(1)
63     compareY = patchWidth*patchWidth * np.random.random((nbits,1))
64     compareY = np.floor(compareY).astype(int)
65
66     return (compareX, compareY)
67
68 def computePixel(img, idx1, idx2, width, center):
69
70     halfWidth = width // 2
71     col1 = idx1 % width - halfWidth
72     row1 = idx1 // width - halfWidth
73     col2 = idx2 % width - halfWidth
74     row2 = idx2 // width - halfWidth
75     return 1 if img[int(center[0]+row1)][int(center[1]+col1)] < img[int(center[0]+row2)][int(center[1]+col2)] else 0
76
77 def computeBrief(img, locs):
78
79     patchWidth = 9
80     nbits = 256
81     compareX, compareY = makeTestPattern(patchWidth, nbits)
82     m, n = img.shape
83
84     halfWidth = patchWidth//2
85
86     locs = np.array(list(filter(lambda x: halfWidth <= x[0] < m-halfWidth and halfWidth <= x[1] < n-halfWidth, locs)))
87     desc = np.array([list(map(lambda x: computePixel(img, x[0], x[1], patchWidth, c), zip(compareX, compareY))) for c in locs])
88
89     return desc, locs
90
91 def corner_detection(img, sigma):
92
93     # fast method
94     result_img = skimage.feature.corner_fast(img, n=PATCHWIDTH, threshold=sigma)
95     locs = skimage.feature.corner_peaks(result_img, min_distance=1)
96     return locs
97
98 def loadVid(path):
99
100     # Create a VideoCapture object and read from input file
101     # If the input is the camera, pass 0 instead of the video file name
102
103     cap = cv2.VideoCapture(path)
104
105     # get fps, width, and height
106     fps = cap.get(cv2.CAP_PROP_FPS)
107     width = cap.get(cv2.CAP_PROP_FRAME_WIDTH)
108     height = cap.get(cv2.CAP_PROP_FRAME_HEIGHT)
109
110     # Append frames to list
111     frames = []
112
113     # Check if camera opened successfully
114     if cap.isOpened() == False:
115         print("Error opening video stream or file")

```

```
116
117     # Read until video is completed
118     while(cap.isOpened()):
119
120         # Capture frame-by-frame
121         ret, frame = cap.read()
122
123         if ret:
124             #Store the resulting frame
125             frames.append(frame)
126         else:
127             break
128
129     # When everything done, release the video capture object
130     cap.release()
131     frames = np.stack(frames)
132
133     return frames, fps, width, height
```

Download data

Download the required data and setup the results directory. If running on colab, DATA\_PARENT\_DIR must be DATA\_PARENT\_DIR = '/content/'

Otherwise, use the local directory of your choosing. Data will be downloaded to DATA\_PARENT\_DIR/hw3\_data and a subdirectory DATA\_PARENT\_DIR/results will be created.

```
1 # Only change this if you are running locally
2 # Default on colab: DATA_PARENT_DIR = '/content/'
3
4 # Data will be downloaded to DATA_PARENT_DIR/hw3_data/
5 # A subdirectory DATA_PARENT_DIR/results will be created
6
7 DATA_PARENT_DIR = '/content/'
8
9 if not os.path.exists(DATA_PARENT_DIR):
10     raise RuntimeError('DATA_PARENT_DIR does not exist: ', DATA_PARENT_DIR)
11
12 RES_DIR = os.path.join(DATA_PARENT_DIR, 'results')
13 if not os.path.exists(RES_DIR):
14     os.mkdir(RES_DIR)
15     print('made directory: ', RES_DIR)
16
17
18 #paths different files are saved to
19 # OPTIONAL:
20 # feel free to change if funning locally
21 ROT_MATCHES_PATH = os.path.join(RES_DIR, 'brief_rot_test.pkl')
22 ROT_INV_MATCHES_PATH = os.path.join(RES_DIR, 'ec_brief_rot_inv_test.pkl')
23 AR_VID_FRAMES_PATH = os.path.join(RES_DIR, 'q_3_1_frames.npy')
24 AR_VID_FRAMES_EC_PATH = os.path.join(RES_DIR, 'q_3_2_frames.npy')
25
26 HW3_SUBDIR = 'hw3_data'
27 DATA_DIR = os.path.join(DATA_PARENT_DIR, HW3_SUBDIR)
28 ZIP_PATH = DATA_DIR + '.zip'
29 if not os.path.exists(DATA_DIR):
30     !wget 'https://www.andrew.cmu.edu/user/hfreeman/data/16720_spring/hw3_data.zip' -O $ZIP_PATH
31     !unzip -qq $ZIP_PATH -d $DATA_PARENT_DIR
```

```
made directory: /content/results
--2024-03-02 16:30:36--  https://www.andrew.cmu.edu/user/hfreeman/data/16720_spring/hw3_data.zip
Resolving www.andrew.cmu.edu (www.andrew.cmu.edu)... 128.2.42.53
Connecting to www.andrew.cmu.edu (www.andrew.cmu.edu)|128.2.42.53|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 36434294 (35M) [application/zip]
Saving to: '/content/hw3_data.zip'
```

```
/content/hw3_data.z 100%[=====] 34.75M  3.21MB/s   in 11s

2024-03-02 16:30:48 (3.06 MB/s) - '/content/hw3_data.zip' saved [36434294/36434294]
```

Q2 Computing Planar Homographies

Q2.1 Feature Detection and Matching

Q2.1.1 (5 points):

How is the FAST detector different from the Harris corner detector that you've seen in the lectures? Can you comment on its computation performance compared to the Harris corner detector?

The FAST detector identifies corners by checking the pixel intensities around a neighbourhood around the particular pixel. If a pixel has a significant number of pixels in its neighbourhood that are either darker or lighter than that pixel, then it is defined as a corner. To make the FAST detector even quicker, a Decision Tree is often used to partition pixels into either of three sets: Darker than Center, Lighter than Center, and Similar to the Center Pixel. The process is iteratively called until there is no more information gain.

The Harris corner detector, on the other hand, identifies corners by analyzing the changes in intensity that occur when a small window around a pixel is shifted in different directions. It computes a corner response function based on the gradients of the image and determines corners where this response is maximized. For this reason, the Harris corner detection is more robust to noise.

The computation performance of the FAST detector is significantly faster than the Harris Corner Detector. This is because the FAST detector only performs comparison of Pixel Intensities, whereas the Harris Corner Detector performs gradient computations. Furthermore, the FAST detector involves around 9 or 12 operations per pixel, this can be further reduced based on some stopping criteria that can effectively say a pixel is not a corner. The FAST detector uses a Decision Tree model that can efficiently partition pixels into corner.

Q2.1.2 (5 points):

How is the BRIEF descriptor different from the filterbanks you've seen in the lectures? Could you use any one of the those filter banks as a descriptor?

The BRIEF descriptor is different from the filterbanks we'd seen earlier in the sense that BRIEF computes the difference in pixel intensities for a random set of pixels based on some sampling strategy. BRIEF is a binary descriptor because it only checks whether the pixel is darker or brighter than the other pixel. In filterbanks, multiple filters are applied on the same image to get different information regarding the image. Filterbanks are more often used as feature extractors rather than feature descriptors.

However, aggregating filterbanks together could result in a descriptor of the image.

Q2.1.3 (5 points):

Describe how the Hamming distance and Nearest Neighbor can be used to match interest points with BRIEF descriptors. What benefits does the Hamming distance have over a more conventional Euclidean distance measure in our setting?

The output of BRIEF is a binary string of 1s and 0s. To help match interest points with BRIEF descriptors, the Euclidean distance may not be suitable. This is because the Euclidean distance operates in a continuous space. This would mean that each BRIEF descriptor would have to be converted into such a space, and then compare the distances. The Hamming distance, however, operates in a binary discrete space, making it more suitable for comparing BRIEF descriptors between two interest points. Using Hamming Distance as the distance function, the Nearest Neighbour for an interest point can be computed based on this distance, effectively matching interest points between two images.

Q2.1.4 (10 points):

Implement the function matchPics()

```
1 def matchPics(I1, I2, ratio, sigma):
2     """
3     Match features across images
4
5     Input
6     -----
7     I1, I2: Source images (RGB or Grayscale uint8)
8     ratio: ratio for BRIEF feature descriptor
9     sigma: threshold for corner detection using FAST feature detector
10
11     Returns
12     -----
13     matches: List of indices of matched features across I1, I2 [p x 2]
14     locs1, locs2: Pixel coordinates of matches [N x 2]
15     """
16
17     # ===== your code here! =====
18
19     # TODO0: Convert images to GrayScale
20     # Input images can be either RGB or Grayscale uint8 (0 -> 255). Both need
21     # to be supported.
22     # Input images must be converted to normalized Grayscale (0.0 -> 1.0)
23     # skimage.color.rgb2gray may be useful if the input is RGB.
24     # TODO0: Detect features in both images
25     # TODO0: Obtain descriptors for the computed feature locations
26     # TODO0: Match features using the descriptors
27     # Convert images to grayscale if they are RGB
28     print("Image1 Shape: "+str(I1.shape))
29     print("Image2 Shape: "+str(I2.shape))
30
31     if len(I1.shape) > 2:
32         I1_gray = skimage.color.rgb2gray(I1)
33     else:
34         I1_gray = I1
35
36     if len(I2.shape) > 2:
37         I2_gray = skimage.color.rgb2gray(I2)
38     else:
39         I2_gray = I2
40
41     print("Image1 Shape: "+str(I1_gray.shape))
42     print("Image2 Shape: "+str(I2_gray.shape))
43
44     # Detect features in both images
45     locs1 = corner_detection(I1_gray, sigma)
46     locs2 = corner_detection(I2_gray, sigma)
47
48     # Compute descriptors for the detected feature locations using BRIEF
49     desc1, locs1 = computeBrief(I1_gray, locs1)
50     desc2, locs2 = computeBrief(I2_gray, locs2)
51
52     # Match features using the computed descriptors
53     matches = briefMatch(desc1, desc2, ratio)
54     # ==== end of code ====
55
56     return matches, locs1, locs2
```

Implement the function displayMatched

```
1 def displayMatched(I1, I2, ratio, sigma):
2     """
3     Displays matches between two images
4
5     Input
6     -----
7     I1, I2: Source images
8     ratio: ratio for BRIEF feature descriptor
9     sigma: threshold for corner detection using FAST feature detector
10     """
11
12     print('Displaying matches for ratio: ', ratio, ' and sigma: ', sigma)
13
14     # ===== your code here! =====
15     # TODO0: Use matchPics and plotMatches to visualize your results
16
17     # Find matches between images
18     matches, locs1, locs2 = matchPics(I1, I2, ratio, sigma)
19     print("Matches: "+str(len(matches)))
20     print("locs1: "+str(len(locs1)))
21     print("locs2: "+str(len(locs2)))
22
23     # Visualize matches
24     plotMatches(I1, I2, matches, locs1, locs2)
25
26     # ==== end of code ====
27
```

Visualize the matches

Use the cell below to visualize the matches. The resulting figure should look similar (but not necessarily identical) to Figure 2.


Feel free to play around with the images and parameters. Please use the original images when submitting the report.

Figure 2 parameters:

- image1\_name = "cv\_cover.jpg"
- image1\_name = "cv\_desk.png"
- ratio = 0.7
- sigma = 0.15

```
1 # Feel free to play around with these parameters
2 # BUT when submitting the report use the original images
3 image1_name = "cv_cover.jpg"
4 image2_name = "cv_desk.png"
5 ratio = 0.7
6 sigma = 0.1
7
8 image1_path = os.path.join(DATA_DIR, image1_name)
9 image2_path = os.path.join(DATA_DIR, image2_name)
10
11 image1 = cv2.imread(image1_path)
12 image2 = cv2.imread(image2_path)
13
14 #bgr to rgb
15 if len(image1.shape) == 3 and image1.shape[2] == 3:
16     image1 = cv2.cvtColor(image1, cv2.COLOR_BGR2RGB)
17
18 if len(image2.shape) == 3 and image2.shape[2] == 3:
19     image2 = cv2.cvtColor(image2, cv2.COLOR_BGR2RGB)
20
21 displayMatched(image1, image2, ratio, sigma)
```

```
Displaying matches for ratio: 0.7 and sigma: 0.1
Image1 Shape: (440, 350, 3)
Image2 Shape: (548, 731, 3)
Image1 Shape: (440, 350)
Image2 Shape: (548, 731)
<ipython-input-10-2fa2590324ba>:75: DeprecationWarning: Conversion of an array with ndim > 0 to a scalar is deprecated, and will error in future. Ensure you extract a single element from your array before performin
    return 1 if img[int(center[0]+row1)][int(center[1]+col1)] < img[int(center[0]+row2)][int(center[1]+col2)] else 0
Matches: 64
locs1: 1813
locs2: 862
```



Q2.1.5 (10 points):

Experiment with different sigma and ratio values. Conduct a small ablation study, and include the figures displaying the matched features with various parameters in your write-up. Explain the effect of these two paremeters respectively.

Here are the results of the ablation study.

1) Varying the Ratio value


We see that by varying the Ratio value from 0.6, 0.7, 0.8, all the way to 0.9, the number of matched interest points increases. This is because the ratio parameter here specifies the maximum ratio between the first best match and second best match between matches of an interest point to two different interest points. By increasing this ratio, we are effectively allowing unstable matches.

2) Varying the Sigma value


We see that by varying the Sigma value from 0.1, 0.15, 0.2, 0.3, all the way to 0.5, the threshold for pixel intensity comparison is altered. This means that as the sigma goes up, less points are marked as interest points as the neighbouring pixels would have to have a higher difference in intensities. This can be seen in the images below

```
1 image1_name = "cv_cover.jpg"
2 image2_name = "cv_desk.png"
3
4 image1_path = os.path.join(DATA_DIR, image1_name)
5 image2_path = os.path.join(DATA_DIR, image2_name)
6
7 image1 = cv2.imread(image1_path)
8 image2 = cv2.imread(image2_path)
9
10 #bgr to rgb
11 if len(image1.shape) == 3 and image1.shape[2] == 3:
12     image1 = cv2.cvtColor(image1, cv2.COLOR_BGR2RGB)
13
14 if len(image2.shape) == 3 and image2.shape[2] == 3:
15     image2 = cv2.cvtColor(image2, cv2.COLOR_BGR2RGB)
16
17 # ===== your code here! =====
18 ratio = [0.6, 0.7, 0.8, 0.9]
19 sigma = [0.1, 0.15, 0.2, 0.3, 0.5]
20
21 for r in ratio:
22     displayMatched(image1, image2, r, 0.15)
23
24 for s in sigma:
25     displayMatched(image1, image2, 0.7, s)
26
27 # ===== end of code =====
```


Displaying matches for ratio: 0.6 and sigma: 0.15  
Image1 Shape: (440, 350, 3)  
Image2 Shape: (548, 731, 3)  
Image1 Shape: (440, 350)  
Image2 Shape: (548, 731)  
<ipython-input-10-2fa2590324ba>:75: DeprecationWarning: Conversion of an array with ndim > 0 to a scalar is deprecated, and will error in future. Ensure you extract a single element from your array before performing  
return 1 if img[int(center[0]+row1)][int(center[1]+col1)] < img[int(center[0]+row2)][int(center[1]+col2)] else 0  
Matches: 10  
locs1: 945  
locs2: 476




Displaying matches for ratio: 0.7 and sigma: 0.15  
Image1 Shape: (440, 350, 3)  
Image2 Shape: (548, 731, 3)  
Image1 Shape: (440, 350)  
Image2 Shape: (548, 731)  
Matches: 24  
locs1: 945  
locs2: 476




Displaying matches for ratio: 0.8 and sigma: 0.15  
Image1 Shape: (440, 350, 3)  
Image2 Shape: (548, 731, 3)  
Image1 Shape: (440, 350)  
Image2 Shape: (548, 731)  
Matches: 61  
locs1: 945  
locs2: 476




Displaying matches for ratio: 0.9 and sigma: 0.15  
Image1 Shape: (440, 350, 3)  
Image2 Shape: (548, 731, 3)  
Image1 Shape: (440, 350)  
Image2 Shape: (548, 731)  
Matches: 125  
locs1: 945  
locs2: 476




Displaying matches for ratio: 0.7 and sigma: 0.1  
Image1 Shape: (440, 350, 3)  
Image2 Shape: (548, 731, 3)  
Image1 Shape: (440, 350)  
Image2 Shape: (548, 731)  
Matches: 64  
locs1: 1813  
locs2: 862




Displaying matches for ratio: 0.7 and sigma: 0.15  
Image1 Shape: (440, 350, 3)  
Image2 Shape: (548, 731, 3)  
Image1 Shape: (440, 350)  
Image2 Shape: (548, 731)  
Matches: 24  
locs1: 945  
locs2: 476




Displaying matches for ratio: 0.7 and sigma: 0.2  
Image1 Shape: (440, 350, 3)  
Image2 Shape: (548, 731, 3)  
Image1 Shape: (440, 350)  
Image2 Shape: (548, 731)  
Matches: 12  
locs1: 601  
locs2: 309




Displaying matches for ratio: 0.7 and sigma: 0.3  
Image1 Shape: (440, 350, 3)  
Image2 Shape: (548, 731, 3)  
Image1 Shape: (440, 350)  
Image2 Shape: (548, 731)  
Matches: 5  
locs1: 320  
locs2: 136





Displaying matches for ratio: 0.7 and sigma: 0.5  
Image1 Shape: (440, 350, 3)  
Image2 Shape: (548, 731, 3)  
Image1 Shape: (440, 350)  
Image2 Shape: (548, 731)  
Matches: 2  
locs1: 154  
locs2: 29



Q2.1.6 (10 points):

Implement the function briefRot

```
1 def briefRot(min_deg, max_deg, deg_inc, ratio, sigma, filename):
2     """
3     Tests Brief with rotations.
4
5     Input
6     -----
7     min_deg: minimum degree to rotate image
8     max_deg: maximum degree to rotate image
9     deg_inc: number of degrees to increment when iterating
10    ratio: ratio for BRIEF feature descriptor
11    sigma: threshold for corner detection using FAST feature detector
12    filename: filename of image to rotate
13
14    """
15
16    if not os.path.exists(RES_DIR):
17        raise RuntimeError('RES_DIR does not exist. did you run all cells?')
18
19    # Read the image and convert bgr to rgb
20    image_path = os.path.join(DATA_DIR, filename)
21    image = cv2.imread(image_path)
22    if len(image.shape) == 3 and image.shape[2] == 3:
23        image = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)
24
25    match_degrees = [] # stores the degrees of rotation
26    match_counts = [] # stores the number of matches at each degree of rotation
27
28    for i in range(min_deg, max_deg, deg_inc):
29        print(i)
30
31        # ===== your code here! =====
32
33        # Rotate Image
34        rotated_image = scipy.ndimage.rotate(image, i, reshape=False)
35
36        matches, locs1, locs2 = matchPics(image, rotated_image, ratio, sigma)
37
38        match_degrees.append(i)
39        match_counts.append(len(matches))
40        # ===== end of code =====
41
42    # Save to pickle file
43    matches_to_save = [match_counts, match_degrees, deg_inc]
44    write_pickle(ROT_MATCHES_PATH, matches_to_save)
45
46 def dispBriefRotHist(matches_path=ROT_MATCHES_PATH):
47     # Check if pickle file exists
48     if not os.path.exists(matches_path):
49         raise RuntimeError('matches_path does not exist. did you call briefRot?')
50
51     # Read from pickle file
52     match_counts, match_degrees, deg_inc = read_pickle(matches_path)
53
54     # Display histogram
55     # Bins are centered and separated every 10 degrees
56     plt.figure()
57     bins = [x - deg_inc/2 for x in match_degrees]
58     bins.append(bins[-1] + deg_inc)
59     plt.hist(match_degrees, bins=bins, weights=match_counts, log=True)
60     #plt.hist(match_degrees, bins=[10 * (x-0.5) for x in range(37)], weights=match_counts, log=True)
61     plt.title("Histogram of BREIF matches")
62     plt.ylabel("# of matches")
63     plt.xlabel("Rotation (deg)")
64     plt.tight_layout()
65
66     output_path = os.path.join(RES_DIR, 'histogram.png')
67     plt.savefig(output_path)
```

Visualize the matches under rotation


See debugging tips in handout.

```
1 # defaults are:
2 # min_deg = 0
3 # max_deg = 360
4 # deg_inc = 10
5 # ratio = 0.7
6 # sigma = 0.15
7 # filename = 'cv_cover.jpg'
8
9 # Controls the rotation degrees
10 min_deg = 0
11 max_deg = 360
12 deg_inc = 10
13
14 # Brief feature descriptor and Fast feature detector paremeters
15 # (change these if you want to use different values)
16 ratio = 0.7
17 sigma = 0.15
18
19 # image to rotate and match
20 # (no need to change this but can if you want to experiment)
21 filename = 'cv_cover.jpg'
22
23 # Call briefRot
24 briefRot(min_deg, max_deg, deg_inc, ratio, sigma, filename)
25
26 Image1 Shape: (440, 350)
27 Image2 Shape: (440, 350)
28
29 20
30 Image1 Shape: (440, 350, 3)
31 Image2 Shape: (440, 350, 3)
32 Image1 Shape: (440, 350)
33 Image2 Shape: (440, 350)
34
35 30
36 Image1 Shape: (440, 350, 3)
37 Image2 Shape: (440, 350, 3)
38 Image1 Shape: (440, 350)
39 Image2 Shape: (440, 350)
40
41 40
42 Image1 Shape: (440, 350, 3)
43 Image2 Shape: (440, 350, 3)
44 Image1 Shape: (440, 350)
45 Image2 Shape: (440, 350)
46
47 50
48 Image1 Shape: (440, 350, 3)
49 Image2 Shape: (440, 350, 3)
50 Image1 Shape: (440, 350)
51 Image2 Shape: (440, 350)
52
53 60
54 Image1 Shape: (440, 350, 3)
55 Image2 Shape: (440, 350, 3)
56 Image1 Shape: (440, 350)
57 Image2 Shape: (440, 350)
58
59 70
60 Image1 Shape: (440, 350, 3)
61 Image2 Shape: (440, 350, 3)
62 Image1 Shape: (440, 350)
63 Image2 Shape: (440, 350)
64
65 80
66 Image1 Shape: (440, 350, 3)
67 Image2 Shape: (440, 350, 3)
68 Image1 Shape: (440, 350)
69 Image2 Shape: (440, 350)
70
71 90
72 Image1 Shape: (440, 350, 3)
73 Image2 Shape: (440, 350, 3)
74 Image1 Shape: (440, 350)
75 Image2 Shape: (440, 350)
76
77 100
78 Image1 Shape: (440, 350, 3)
79 Image2 Shape: (440, 350, 3)
80 Image1 Shape: (440, 350)
81 Image2 Shape: (440, 350)
82
83 110
84 Image1 Shape: (440, 350, 3)
85 Image2 Shape: (440, 350, 3)
86 Image1 Shape: (440, 350)
87 Image2 Shape: (440, 350)
88
89 120
90 Image1 Shape: (440, 350, 3)
91 Image2 Shape: (440, 350, 3)
92 Image1 Shape: (440, 350)
93 Image2 Shape: (440, 350)
94
95 130
96 Image1 Shape: (440, 350, 3)
97 Image2 Shape: (440, 350, 3)
98 Image1 Shape: (440, 350)
99 Image2 Shape: (440, 350)
```

Plot the histogram

See debugging tips in handout.

```
1 dispBriefRotHist()
```



Explain why you think the BRIEF descriptor behaves this way: BRIEF is not rotation invariant. The binary descriptor for the matched points will change as the orientation of the image changes. This reduces the number of matched points. This can be seen in the histogram. For small rotations, like 0, 10, and 350, we see that there are still a large no. of matches. This is probably due to the fact that such small rotations would not change the final binary descriptor too much, which depends on a certain sampling strategy.

Q2.1.7.1 (Extra Credit - 5 points):

Design a fix to make BRIEF more rotation invariant. Feel free to make any helper functions as necessary. But you cannot use any additional OpenCV or Scikit-Image functions.



```

1 # ===== your code here! =====
2 # TODO: Define any helper functions here
3 # (Feel free to put anything in its own cell)
4
5 # TODO: Feel free to modify the inputs and the function body as necessary
6 # This is only an outline
7 def briefRotInvEc(min_deg, max_deg, deg_inc, ratio, sigma, filename):
8     """
9     Rotation invariant Brief.
10
11     Input
12     -----
13     min_deg: minimum degree to rotate image
14     max_deg: maximum degree to rotate image
15     deg_inc: number of degrees to increment when iterating
16     ratio: ratio for BRIEF feature descriptor
17     sigma: threshold for corner detection using FAST feature detector
18     filename: filename of image to rotate
19
20     """
21
22     if not os.path.exists(RES_DIR):
23         raise RuntimeError('RES_DIR does not exist. did you run all cells?')
24
25     #Read the image and convert bgr to rgb
26     image_path = os.path.join(DATA_DIR, filename)
27     image = cv2.imread(image_path)
28     if len(image.shape) == 3 and image.shape[2] == 3:
29         image = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)
30
31     match_degrees = [] # stores the degrees of rotation
32     match_counts = [] # stores the number of matches at each degree of rotation
33
34     for i in range(min_deg, max_deg, deg_inc):
35         print(i)
36
37         # TODO: Rotate Image (Hint: use scipy.ndimage.rotate)
38
39         # TODO: Brief matcher that is rotation invariant
40         # Feel free to define additional helper functions as necessary
41
42         # TODO: visualizes matches at at least 3 different orientations
43         # to include in your report
44         # (Hint: use plotMatches)
45
46         # TODO: Update match_degrees and match_counts (see descriptions above)
47
48     # Save to pickle file
49     matches_to_save = [match_counts, match_degrees, deg_inc]
50     write_pickle(ROT_INV_MATCHES_PATH, matches_to_save)
51
52 # ===== end of code =====

```

Visualize your implemented function

```

1 min_deg = 0
2 max_deg = 360
3 deg_inc = 10
4 filename = 'cv_cover.jpg'
5
6 # ===== your code here! =====
7 # TODO: Call briefRotInvEc and visualize
8
9 # ===== end of code =====
10

```

Plot Histogram

```

1 dispBriefRotHist(matches_path=ROT_INV_MATCHES_PATH)

```

```

-----
RuntimeError                                Traceback (most recent call last)
<ipython-input-24-8890031f161e> in <cell line: 1>()
----> 1 dispBriefRotHist(matches_path=ROT_INV_MATCHES_PATH)

<ipython-input-19-b92561582089> in dispBriefRotHist(matches_path)
    47     # Check if pickle file exists
    48     if not os.path.exists(matches_path):
----> 49         raise RuntimeError('matches_path does not exist. did you call briefRot?')
    50
    51     # Read from pickle file

RuntimeError: matches_path does not exist. did you call briefRot?

```

---

Compare the histograms with an without rotation invariance. Explain your rotation invariant design and how you selected any parameters that you used: YOUR ANSWER HERE...

---

Q2.1.7.2 (Extra Credit - 5 points):

Design a fix to make BRIEF more scale invariant. Feel free to make any helper functions as necessary. But you cannot use any additional OpenCV or Scikit-Image functions.

```
1 # ===== your code here! =====
2 # TODO: Define any helper functions here
3 # (Feel free to put anything in its own cell)
4
5 # TODO: Modify the inputs and the function body as necessary
6 def briefScaleInvEc(ratio, sigma, filename):
7
8     #Read the image and convert bgr to rgb
9     image_path = os.path.join(DATA_DIR, filename)
10    image = cv2.imread(image_path)
11    if len(image.shape) == 3 and image.shape[2] == 3:
12        image = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)
13
14    match_scales = [] # stores the scaling factors
15    match_counts = [] # stores the number of matches at each scaling factor
16
17    for i in [1]:
18        # Scale Image
19        image_scale = cv2.resize(image,(int(image.shape[1]/(2**i)),
20                                int(image.shape[0]/(2**i))),
21                                interpolation = cv2.INTER_AREA)
22
23        # TODO: Brief matcher that is scale invariant
24        # Feel free to define additional helper functions as necessary
25
26        # Compare to regular matchPics
27        matches_orig, locs1_orig, locs2_orig = matchPics(image,
28                                                         image_scale,
29                                                         ratio, sigma)
30
31        print('plotting non-scale invariant scale: ', 2**i)
32        plotMatches(image, image_scale, matches_orig, locs1_orig,
33                   locs2_orig)
34        print('plotting scale-invariant: ', 2**i)
35        plotMatches(image, image_scale, matches, locs1, locs2)
36
37 # ===== end of code =====
```

Visualize your implemented function

```
1 # ===== your code here! =====
2 # TODO: Call briefScaleInvEc and visualize
3 # You may change any parameters and the function body as necessary
4
5 filename = 'cv_cover.jpg'
6
7 ratio = 0.7
8 sigma = 0.15
9
10 briefScaleInvEc(ratio, sigma, filename)
11 # ===== end of code =====
```

Explain your scale invariant design and how you selected any parameters that you used: YOUR ANSWER HERE...

## Q2.2 Homography Computation

### Q2.2.1 (15 Points): Compute H

Implement the function computeH

```
1 def computeH(x1, x2):
2     """
3     Compute the homography between two sets of points
4
5     Input
6     -----
7     x1, x2: Sets of points
8
9     Returns
10    -----
11    H2to1: 3x3 homography matrix that best transforms x2 to x1
12    """
13
14    if x1.shape != x2.shape:
15        raise RuntimeError('number of points do not match')
16
17    # ===== your code here! =====
18    # TODO: Compute the homography between two sets of points
19
20    # Construct matrix A
21
22    N = x1.shape[0]
23    A = np.zeros((2 * N, 9))
24
25    for i in range(N):
26        A[2 * i] = [x2[i, 0], x2[i, 1], 1, 0, 0, 0, -x2[i, 0] * x1[i, 0], -x2[i, 1] * x1[i, 0], -x1[i, 0]]
27        A[2 * i + 1] = [0, 0, 0, x2[i, 0], x2[i, 1], 1, -x2[i, 0] * x1[i, 1], -x2[i, 1] * x1[i, 1], -x1[i, 1]]
28
29    # Perform SVD on A
30    _, _, V = np.linalg.svd(A)
31
32    # Extract homography matrix H2to1 from the last column of V
33    H2to1 = V[-1].reshape(3, 3)
34
35    return H2to1
```

### Q2.2.2 (10 points): ComputeH\_norm

Implement the function computeH\_norm

```
1 def computeH_norm(x1, x2):
2     """
3     Compute the homography between two sets of points using normalization
4
5     Input
6     -----
7     x1, x2: Sets of points
8
9     Returns
```

```

11 H2to1: 3x3 homography matrix that best transforms x2 to x1
12 """
13
14 # ===== your code here! =====
15
16 # TODO: Compute the centroid of the points
17 centroid1 = np.mean(x1, axis=0)
18 centroid2 = np.mean(x2, axis=0)
19
20 # TODO: Shift the origin of the points to the centroid
21 x1_centered = x1 - centroid1
22 x2_centered = x2 - centroid2
23
24 # TODO: Normalize the points so that the largest distance from the
25 # origin is equal to sqrt(2)
26 max_dist = np.sqrt(2)
27 scale1 = max_dist / np.max(np.linalg.norm(x1_centered, axis=1))
28 scale2 = max_dist / np.max(np.linalg.norm(x2_centered, axis=1))
29
30
31 # TODO: Similarity transform 1
32 T1 = np.array([[scale1, 0, -scale1 * centroid1[0]],
33               [0, scale1, -scale1 * centroid1[1]],
34               [0, 0, 1]])
35
36 # TODO: Similarity transform 2
37 T2 = np.array([[scale2, 0, -scale2 * centroid2[0]],
38               [0, scale2, -scale2 * centroid2[1]],
39               [0, 0, 1]])
40
41 x1_centered = np.column_stack((x1_centered, np.ones(x1_centered.shape[0])))
42 x2_centered = np.column_stack((x2_centered, np.ones(x2_centered.shape[0])))
43
44 x1_norm = T1 @ x1_centered.T
45 x2_norm = T2 @ x2_centered.T
46 # TODO: Compute homography
47 H2to1 = computeH(x1, x2)
48
49 # TODO: Denormalization\
50 #H2to1 = np.linalg.inv(T1) @ H2to1 @ T2
51 H2to1 /= H2to1[2, 2]
52
53
54 # ==== end of code ====
55
56 return H2to1

```

Q2.2.3 (25 points): ComputeH\_ransac

Implement RANSAC

```

1 def computeH_ransac(locs1, locs2, max_iters, inlier_tol):
2     """
3     Estimate the homography between two sets of points using ransac
4
5     Input
6     -----
7     locs1, locs2: Lists of points
8     max_iters: the number of iterations to run RANSAC for
9     inlier_tol: the tolerance value for considering a point to be an inlier
10
11     Returns
12     -----
13     bestH2to1: 3x3 homography matrix that best transforms locs2 to locs1
14     inliers: indices of RANSAC inliers
15
16     """
17
18     # ===== your code here! =====
19
20     # TODO:
21     # Compute the best fitting homography using RANSAC
22     # given a list of matching points locs1 and loc2
23     bestH2to1 = None
24     best_inliers = []
25
26     flag=False
27     print("Max Iterations: "+str(max_iters))
28     for m in range(max_iters):
29
30         locs1_swapped = np.column_stack((locs1[:, 1], locs1[:, 0]))
31         locs2_swapped = np.column_stack((locs2[:, 1], locs2[:, 0]))
32
33         rand_indices = np.random.choice(len(matches), 4, replace=False)
34         rand_locs1 = [locs1_swapped[i] for i in rand_indices]
35         rand_locs2 = [locs2_swapped[i] for i in rand_indices]
36
37         rand_locs1=np.array(rand_locs1)
38         rand_locs2=np.array(rand_locs2)
39
40         H = computeH_norm(rand_locs1, rand_locs2)
41         #print(H)
42         #bestH2to1= H
43
44
45         inliers = []
46         for i in range(len(matches)):
47             p1 = np.append(locs1_swapped[i], 1)
48             p2 = np.append(locs2_swapped[i], 1)
49             p2_transformed = H @ p2
50             print(np.linalg.norm(p1 - p2_transformed))
51             if np.linalg.norm(p1 - p2_transformed) < inlier_tol:
52                 print("inlier!!")
53                 inliers.append(i)
54
55         if len(inliers) > len(best_inliers):
56             flag=True
57             print("good H found")
58             best_inliers = inliers
59             bestH2to1 = H
60
61 if flag==False:
62     bestH2to1 = H

```

```
64 # ==== end of code ====
65
66 return bestH2to1, best_inliers
```

Q2.2.4 (10 points): compositeH

Implement the function compositeH

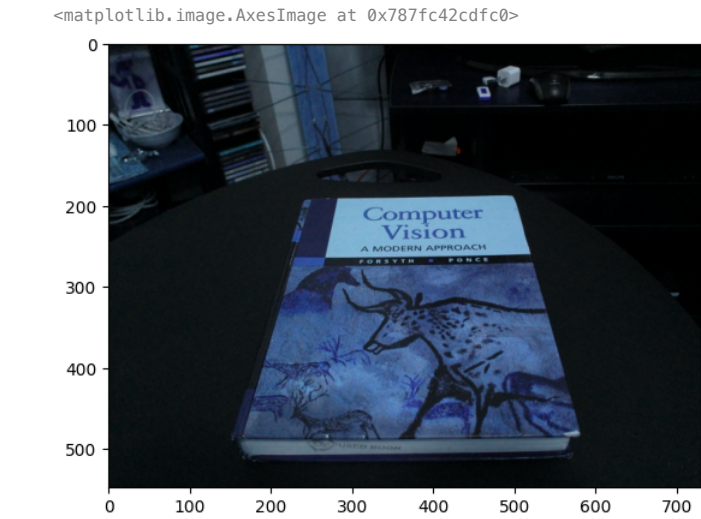
```
1 def compositeH(H2to1, template, img):
2     """
3     Returns the composite image.
4
5     Input
6     -----
7     H2to1: Homography from image to template
8     template: template image to be warped
9     img: background image
10
11     Returns
12     -----
13     composite_img: Composite image
14
15     """
16
17     # ===== your code here! =====
18     # TODO: Create a composite image after warping the template image on top
19     # of the image using the homography
20
21     h, w, _ = img.shape
22     plt.imshow(template)
23     warped_template = cv2.warpPerspective(template, H2to1, (w, h))
24     composite_img = np.where(warped_template != 0, warped_template, img)
25
26     # ==== end of code ====
27
28     return warped_template
```

Implement the function warpImage

```
1 hp_cover = skimage.io.imread(os.path.join(DATA_DIR, 'hp_cover.jpg'))
2 cv_cover = skimage.io.imread(os.path.join(DATA_DIR, 'cv_cover.jpg'))
3 cv_desk = skimage.io.imread(os.path.join(DATA_DIR, 'cv_desk.png'))
4 cv_desk = cv_desk[:, :, :3]
```

```
1 image1 = cv_cover
2 image2 = cv_desk
3 image3= hp_cover
4
5 #bgr to rgb
6 if len(image1.shape) == 3 and image1.shape[2] == 3:
7     image1 = cv2.cvtColor(image1, cv2.COLOR_BGR2RGB)
8
9 if len(image2.shape) == 3 and image2.shape[2] == 3:
10     image2 = cv2.cvtColor(image2, cv2.COLOR_BGR2RGB)
```

```
1 plt.imshow(image2)
```



```
1 matches, locs1, locs2 = matchPics(cv_cover, cv_desk, ratio, sigma)
```

```
Image1 Shape: (440, 350)
Image2 Shape: (548, 731, 3)
Image1 Shape: (440, 350)
Image2 Shape: (548, 731)
<ipython-input-10-2fa2590324ba>:75: DeprecationWarning: Conversion of an array with ndim > 0 to a scalar is deprecated, and will error in future. Ensure you extract a single element from your array before performing scalar operations
return 1 if img[int(center[0]+row1)][int(center[1]+col1)] < img[int(center[0]+row2)][int(center[1]+col2)] else 0
```

```
1 matched_locs1 = locs1[matches[:, 0]] #y,x
2
3 scaling_factor = [hp_cover.shape[0] / cv_cover.shape[0], hp_cover.shape[1] / cv_cover.shape[1]] #y,x
4 print(scaling_factor)
5 scaled_locs1 = np.round(matched_locs1 * scaling_factor).astype(int) #y,x
6
7 matched_locs2 =locs2[matches[:,1]]


[0.6704545454545454, 0.5714285714285714]
```

```
1 locs1_swapped = np.column_stack((scaled_locs1[:, 1], scaled_locs1[:, 0]))
2 locs2_swapped = np.column_stack((matched_locs2[:, 1], matched_locs2[:, 0]))
```

```
1 rand_indices = np.random.choice(len(matches), 4, replace=False)
2
3 rand_locs1 = [locs1_swapped[i] for i in rand_indices]
4 rand_locs2 = [locs2_swapped[i] for i in rand_indices]
```

```
1 rand_locs1=np.array(rand_locs1)
2 rand_locs2=np.array(rand_locs2)
```

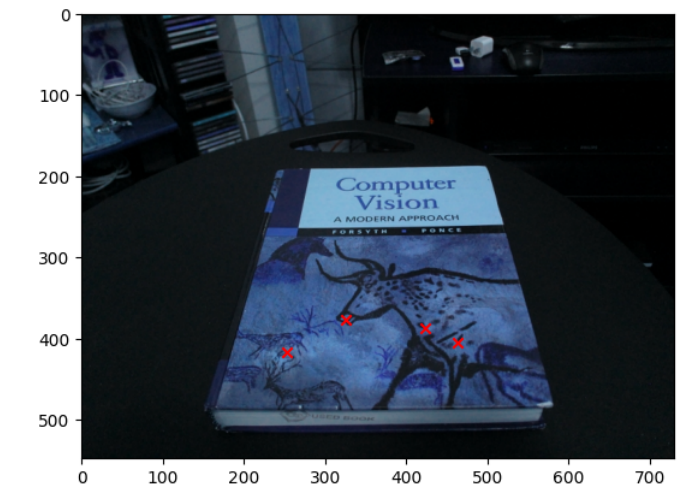
```
1 plt.imshow(image1)
2 plt.scatter(rand_locs1[:, 0], rand_locs1[:, 1], c='r', marker='x')
3 plt.show()
```



```
1 rand_locs2[:, 0]

array([253, 423, 326, 464])

1 plt.imshow(image2)
2 plt.scatter(rand_locs2[:, 0], rand_locs2[:, 1], c='r', marker='x')
3 plt.show()
```



```
1 x1

array([[ 43, 248],
       [130, 226],
       [ 79, 217],
       [150, 241]])

1 x1 = rand_locs1
2 x2 = rand_locs2
3
4 centroid1 = np.mean(x1, axis=0)
5 centroid2 = np.mean(x2, axis=0)
6 print(centroid1)
7 print(centroid2)

[100.5 233. ]
[366.5 396.25]

1 x1_centered = x1 - centroid1
2 x2_centered = x2 - centroid2

1 max_dist = np.sqrt(2)
2 scale1 = max_dist / np.max(np.linalg.norm(x1_centered, axis=1))
3 scale2 = max_dist / np.max(np.linalg.norm(x2_centered, axis=1))

1 # TODO: Similarity transform 1
2 T1 = np.array([[scale1, 0, -scale1 * centroid1[0]],
3               [0, scale1, -scale1 * centroid1[1]],
4               [0, 0, 1]])
5
6 # TODO: Similarity transform 2
7 T2 = np.array([[scale2, 0, -scale2 * centroid2[0]],
8               [0, scale2, -scale2 * centroid2[1]],
9               [0, 0, 1]])

1 x1_centered = np.column_stack((x1, np.ones(x1_centered.shape[0])))
2 x2_centered = np.column_stack((x2, np.ones(x2_centered.shape[0])))

1 x1_centered

array([[ 43., 248.,  1.],
       [130., 226.,  1.],
       [ 79., 217.,  1.],
       [150., 241.,  1.]])

1 x1_norm = T1 @ x1_centered.T
2 print(x1_norm)

[[-1.36841747  0.70205766 -0.51166914  1.17802895]
 [ 0.35697847 -0.16658995 -0.38077704  0.19038852]
 [ 1.         1.         1.         1.        ]]


1 x2_centered

array([[253., 417.,  1.],
       [423., 387.,  1.],
       [326., 376.,  1.],
       [464., 405.,  1.]])


1 x2_norm = T2 @ x2_centered.T
2 print(x2_norm)

[[-1.39115638  0.69251397 -0.49640382  1.19504623]
 [ 0.25433035 -0.11337618 -0.24820191  0.10724774]
 [ 1.         1.         1.         1.        ]]
```

```
1 plt.imshow(image1)
2 plt.scatter(x1_norm[:, 0], x1_norm[:, 1], c='r', marker='x')
3 plt.show()
```



```
1 plt.imshow(image2)
2 plt.scatter(x2_norm[:, 0], x2_norm[:, 1], c='r', marker='x')
3 plt.show()
```



```
1 H2to1 = computeH(x1, x2)
2 H2to1 /= H2to1[2, 2]
```

```
1 H2to1
array([[ 9.44905656e-01,  1.66259041e-01, -2.13384591e+02],
       [-2.32861805e-01,  2.42027131e+00, -4.02394287e+02],
       [-1.03264707e-03,  3.52689352e-03,  1.00000000e+00]])
```


```
1 H2to1 = np.linalg.inv(T1) @ H2to1 @ T2
2 H2to1 /= H2to1[2, 2]
```

```
1 H2to1
array([[ 3.19066558e+00, -7.75134174e-01,  1.58908963e+04],
       [ 5.44528878e-01,  5.86821347e-01,  1.51848292e+04],
       [ 1.70581448e-05, -3.56233972e-05,  1.00000000e+00]])
```

```
1 height, width = image2.shape[:2]
2 warped_image2 = cv2.warpPerspective(image1, np.linalg.inv(H2to1), (width, height))
```

```
1 plt.imshow(warped_image2)
```


<matplotlib.image.AxesImage at 0x787fc3adb4f0>



```
1 composite_img = np.where(warped_image2 != 0, warped_image2, image2.mean(2))
```

```
1 plt.imshow(composite_img)
```

<matplotlib.image.AxesImage at 0x787fc39a9b40>



```

1 def warpImage(ratio, sigma, max_iters, inlier_tol):
2     """
3     Warps hp_cover.jpg onto the book cover in cv_desk.png.
4
5     Input
6     -----
7     ratio: ratio for BRIEF feature descriptor
8     sigma: threshold for corner detection using FAST feature detector
9     max_iters: the number of iterations to run RANSAC for
10    inlier_tol: the tolerance value for considering a point to be an inlier
11
12    """
13
14    hp_cover = skimage.io.imread(os.path.join(DATA_DIR, 'hp_cover.jpg'))
15    cv_cover = skimage.io.imread(os.path.join(DATA_DIR, 'cv_cover.jpg'))
16    cv_desk = skimage.io.imread(os.path.join(DATA_DIR, 'cv_desk.png'))
17    cv_desk = cv_desk[:, :, :3]
18
19    # ===== your code here! =====
20
21    # TODO: match features between cv_desk and cv_cover using matchPics
22    matches, locs1, locs2 = matchPics(cv_cover, cv_desk, ratio, sigma)
23
24    # TODO: Scale matched pixels in cv_cover to size of hp_cover
25    matched_locs1 = locs1[matches[:, 0]]
26    scaling_factor = [hp_cover.shape[0] / cv_cover.shape[0], hp_cover.shape[1] / cv_cover.shape[1]]
27    scaled_locs1 = np.round(matched_locs1 * scaling_factor).astype(int)
28    matched_locs2 = locs2[matches[:, 1]]
29
30    # TODO: Get homography by RANSAC using computeH_ransac
31    H, _ = computeH_ransac(matched_locs1, matched_locs2, max_iters, inlier_tol)
32    print("Compute H ransac done")
33    print(H.dtype)
34
35    # TODO: Overlay using compositeH to return composite_img
36    composite_img = compositeH(H, hp_cover, cv_desk)
37
38    # ===== end of code =====
39
40    plt.imshow(composite_img)
41    plt.show()

```

Visualize composite image

```

1 # defaults are:
2 # ratio = 0.7
3 # sigma = 0.15
4 # max_iters = 600
5 # inlier_tol = 1.0
6
7 # (no need to change this but can if you want to experiment)
8 ratio = 0.7
9 sigma = 0.15
10 max_iters = 600
11 inlier_tol = 50.0
12
13 warpImage(ratio, sigma, max_iters, inlier_tol)

```

Q2.2.5 (10 points):

Conduct ablation study with various max\_iters and inlier\_tol values. Plot the result images and explain the effect of these two parameters respectively.

```

1 # ===== your code here! =====
2 # Experiment with different max_iters and inlier_tol values.
3 # Include the result images in the write-up.
4
5 # ===== end of code =====

```

---


Explain the effect of max\_iters and inlier\_tol: YOUR ANSWER HERE...

---


Q3 Create a Simple Panorama




**Original (Sigma=0.15, Ratio=0.7)**




**(Sigma=0.15, Ratio=0.6)**



**(Sigma=0.15, Ratio=0.8)**




**(Sigma=0.15, Ratio=0.9)**







**Original (Sigma=0.15, Ratio=0.7)**




**(Sigma=0.1, Ratio=0.7)**




**(Sigma=0.2, Ratio=0.7)**



**(Sigma=0.3, Ratio=0.7)**



(Sigma=0.5, Ratio=0.7)



$$1.201) \quad u_1 = K_1 [I \ 0] X$$

$$u_2 = K_2 [R \ 0] X$$

$$X = K_1^{-1} u_1.$$

$$\text{So, } u_2 = K_2 [R \ 0] K_1^{-1} u_1$$

$$\underline{\underline{u_2 = K_2 R K_1^{-1} u_1}}$$

$$\underline{\underline{u_2 = H u_1}}$$



Let's define  $h$  as  $\begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix}$

Let's use the points  $(x_1^i, y_1^i)$  and  $(x_2^i, y_2^i)$  such that  $x_2^i \equiv H \cdot x_1^i$  ( $i \in \{1, \dots, N\}$ )

To compute  $h$ ,

$$\lambda \begin{bmatrix} x_2^i \\ y_2^i \\ 1 \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix} \begin{bmatrix} x_1^i \\ y_1^i \\ 1 \end{bmatrix}$$

$$\begin{aligned} \lambda \overbrace{(gx_1^i + hy_1^i + i)}^{x_2^i} &= ax_1^i + by_1^i + c \\ \lambda (gx_1^i + hy_1^i + i) &= dx_1^i + ey_1^i + f \end{aligned}$$

These can be written as:

$$ax_1^i + by_1^i + c - gx_1^i x_2^i - hy_1^i x_2^i - i x_2^i = 0$$

$$dx_1^i + ey_1^i + f - gx_1^i y_2^i - hy_1^i y_2^i - i y_2^i = 0$$

$$A_i = \begin{bmatrix} x_1^i & y_1^i & 1 & 0 & 0 & 0 & -x_1^i x_2^i & -x_2^i y_1^i & -x_2^i \\ 0 & 0 & 0 & x_1^i & y_1^i & 1 & -x_1^i y_2^i & -y_1^i y_2^i & -y_2^i \end{bmatrix}$$